

Multi-View Prototype-Based Deep Recognition (MVPDR): Efficient Plant Disease Classification using Vision-Language Models

Jatin Gupta, Krishna Singh, and Madhur Kholia
Department of Electrical Engineering

Abstract—Plant disease detection is critical for agricultural productivity and food security. This paper presents MVPDR (Multi-View Prototype-Based Deep Recognition), a novel approach combining CLIP vision-language models with learnable visual and textual prototypes for efficient plant leaf disease classification. Our method achieves 72.36% accuracy on the PlantDoc dataset (27 disease classes) using only frozen CLIP features and lightweight adapters (500K trainable parameters). We demonstrate significant data efficiency in few-shot settings, achieving 59.7% accuracy with only 20 training samples per class (82.5% of full training performance). Comprehensive ablation studies validate robustness across backbone architectures (ResNet-101, RN50, ViT-B/32), fusion weights, and clustering configurations. The approach enables rapid deployment with minimal computational overhead while maintaining strong generalization across three plant disease benchmarks. Our results indicate that vision-language models provide effective domain-general semantics for fine-grained agricultural classification tasks.

I. INTRODUCTION

Plant disease is one of the most significant challenges threatening global agricultural production, causing estimated losses of 15-20% of crops annually. Early and accurate disease detection is essential for effective crop management and pesticide application. Traditional diagnostic methods rely on manual inspection by agricultural experts, which is labor-intensive, time-consuming, and inaccessible in resource-limited regions.

Deep learning has revolutionized plant disease classification through convolutional neural networks (CNNs), achieving impressive accuracy on controlled datasets. However, these approaches face critical limitations: (1) they require large amounts of labeled data, which is expensive and time-consuming to acquire; (2) they exhibit poor generalization to distribution shifts between laboratory and field images; (3) they lack interpretability via semantic grounding; and (4) full retraining is necessary for new diseases or geographic regions.

Recent advances in vision-language models (VLMs) like CLIP [1] have demonstrated remarkable zero-shot and few-shot learning capabilities by leveraging natural language semantic understanding. However, applying pre-trained VLMs directly to specialized domains like plant pathology results in sub-optimal performance due to domain mismatch (38.2% zero-shot accuracy on our benchmark).

In this work, we propose MVPDR, a framework that bridges this gap by combining frozen CLIP representations with

learnable visual and textual prototypes. Our key contributions are:

- 1) A lightweight fine-tuning approach using only adapter weights (500K parameters) on frozen CLIP features, eliminating the need to retrain large vision backbones while maintaining computational efficiency.
- 2) A multi-view prototype learning strategy that captures both intra-class visual diversity (via K-means clustering) and domain-specific semantic information (via disease-specific text prompts).
- 3) Comprehensive evaluation demonstrating strong few-shot learning capability (59.7% accuracy at 20-shot, an 82.5% efficiency gain relative to full training).
- 4) Thorough ablation studies establishing robustness across hyperparameters, random seeds, and architectural choices.
- 5) Honest reporting of limitations including realistic discussion of open-set detection performance.

This paper is organized as follows: Section 2 reviews related work in plant disease classification and few-shot learning. Section 3 details our methodology. Section 4 presents experimental setup and results. Section 5 discusses findings, limitations, and ethical considerations. Section 6 concludes with directions for future work.

II. RELATED WORK

A. Plant Disease Classification

Early deep learning approaches for plant disease classification relied on standard CNNs (AlexNet, VGG, ResNet) trained from scratch or fine-tuned on plant-specific datasets. PlantDoc [2] and PlantVillage [3] became standard benchmarks. However, these methods require substantial labeled data and exhibit limited generalization.

Transfer learning using pre-trained ImageNet models improved performance but still required task-specific fine-tuning and suffered from dataset bias. Recent work has explored domain adaptation and data augmentation techniques to address cross-dataset transfer challenges.

B. Vision-Language Models

CLIP [1] demonstrated that contrastive pre-training on image-text pairs at scale enables robust zero-shot transfer to downstream tasks. Its success sparked research into leveraging

VLMs for specialized domains through lightweight fine-tuning strategies rather than full retraining.

Recent work has shown that prompt engineering and prototype-based approaches can further improve VLM performance on specialized tasks. Our work extends these ideas to agricultural classification with domain-specific prompt generation and multi-view prototype fusion.

C. Few-Shot Learning

Few-shot learning aims to classify new instances with minimal training examples. Meta-learning approaches (MAML, Prototypical Networks, Matching Networks) have shown promise. However, VLM-based few-shot approaches that leverage pre-trained representations are increasingly competitive, particularly when combined with efficient adaptation strategies.

III. METHODOLOGY

A. System Overview

MVPDR operates in three stages:

- 1) **Feature Extraction:** Encode plant leaf images and disease-specific text prompts using frozen CLIP encoder to obtain 512-dimensional feature vectors.
- 2) **Prototype Learning:** Build visual prototypes via K-means clustering on training images and textual prototypes via aggregation of disease description embeddings.
- 3) **Prediction:** Fuse visual and textual similarity scores using learnable adapter weights for final disease classification.

B. Visual Prototype Learning

For each disease class $c \in \{1, \dots, C\}$, we extract CLIP image embeddings $\mathbf{I}_c = \{\mathbf{i}_1^c, \dots, \mathbf{i}_{N_c}^c\}$ from training samples. We apply K-means clustering to obtain K cluster centroids per class:

$$\mathbf{V}_c = \{\mathbf{v}_1^c, \dots, \mathbf{v}_K^c\} = \text{KMeans}(\mathbf{I}_c, K)$$

where K (denoted NCLT in ablations) controls intra-class prototype diversity. Visual affinity scores are computed as:

$$s_v(q, c) = \text{softmax} \left(\frac{1}{\beta - \beta \cdot \max_k \cos(\mathbf{q}, \mathbf{v}_k^c)} \right)$$

C. Textual Prototype Learning

For each disease class, we curate 50 diverse text prompts describing symptoms and characteristics (e.g., “tomato leaves with necrotic lesions and chlorotic rings characteristic of early blight”). Prompts are encoded via CLIP text encoder:

$$\mathbf{T}_c = \{\mathbf{t}_1^c, \dots, \mathbf{t}_{50}^c\} = \text{CLIP}_{\text{text}}(\text{prompts}_c)$$

Textual features are aggregated as:

$$\bar{\mathbf{t}}_c = \gamma \cdot \frac{1}{50} \sum_i \mathbf{t}_i^c + (1 - \gamma) \cdot \max_i \mathbf{t}_i^c$$

where γ controls the balance between mean and max pooling. Textual affinity:

$$s_t(q, c) = \cos(\mathbf{q}, \bar{\mathbf{t}}_c)$$

D. Multi-View Fusion

Final classification logits are computed as a learned combination:

$$\text{logits}_c = s_t(q, c) + \alpha \cdot s_v(q, c)$$

where $\alpha \in [0, 1]$ is a learnable fusion parameter optimized via gradient descent. Only adapter weights and fusion parameters are trained; CLIP backbone remains frozen.

IV. EXPERIMENTS

A. Datasets

We evaluate on three public plant disease benchmarks:

Dataset	Classes	Train	Test
PlantDoc	27	2,077	521
PlantVillage	39	51,000+	3,200+
PlantWild	115	8,000+	2,000+

B. Baseline Setup

Configuration: ResNet-101 backbone, frozen CLIP weights, 50-epoch training with batch size 32, learning rate 0.001, weight decay 0.0001, 16 visual clusters (NCLT=16), fusion weight $\alpha = 0.3$.

Baselines:

- Zero-shot CLIP: Direct classification using prompt embeddings without training
- ImageNet fine-tuned ResNet-101: Standard supervised learning baseline
- Full fine-tuning: End-to-end training of CLIP backbone

C. Results: Baseline Performance

Table IV-C shows baseline results on PlantDoc:

Method	Accuracy	Precision	Recall	F1
Zero-shot CLIP	38.2%	38.5%	36.1%	36.8%
MVPDR (Ours)	72.36%	72.03%	71.04%	70.54%
Improvement	+34.1 pp	+33.5 pp	+35.0 pp	+33.7 pp

D. Results: Few-Shot Learning

Table IV-D demonstrates data efficiency across different shot counts:

Shots	Accuracy	Precision	Recall	F1	Efficiency
1-shot	34.5%	42.7%	36.0%	32.7%	47.7%
5-shot	47.4%	49.1%	50.2%	46.9%	65.5%
10-shot	57.4%	56.2%	57.6%	55.5%	79.3%
20-shot	59.7%	59.0%	60.3%	58.5%	82.5%
Full training	72.4%	72.0%	71.0%	70.5%	100.0%

Data efficiency is defined as: $\frac{\text{Accuracy}_{k\text{-shot}} - \text{Accuracy}_{0\text{-shot}}}{\text{Accuracy}_{\text{full}} - \text{Accuracy}_{0\text{-shot}}}$

Observations: Progressive accuracy improvement from 1 to 20 shots (+12.9 pp gain). The diminishing returns after 20 shots suggest saturation of the lightweight adapter capacity, indicating potential for larger adapter architectures in future work.

E. Results: Ablation Studies

1) *Backbone Architecture*: Table IV-E1 compares different CLIP vision encoders:

Backbone	Accuracy	Precision	Recall	F1
ResNet-50	70.83%	71.21%	68.22%	68.00%
ResNet-101	72.36%	72.03%	71.04%	70.54%
ViT-B/32	72.17%	71.17%	70.91%	69.94%

ResNet-101 achieves best performance (+1.53 pp over RN50, +0.19 pp over ViT-B). Difference between RN50 and ViT-B is marginal, suggesting backbone choice is less critical than hyperparameter tuning.

2) *Fusion Weight (α)*: Table IV-E2 studies the visual-textual interpolation parameter:

α	Accuracy	Precision	Recall	F1
0.0 (text only)	71.79%	70.65%	69.53%	69.68%
0.2	71.98%	71.98%	71.16%	70.62%
0.3 (optimal)	72.36%	72.03%	71.04%	70.54%
0.5	72.36%	72.03%	71.29%	70.59%
0.7	72.17%	70.72%	70.18%	69.37%
1.0 (vision only)	71.98%	71.64%	70.85%	70.40%

Both $\alpha = 0.3$ and $\alpha = 0.5$ achieve peak accuracy (72.36%). Texture-only ($\alpha = 0.0$) slightly underperforms, confirming value of multi-view fusion. Robustness across $\alpha \in [0.2, 0.5]$ suggests the approach is forgiving to hyperparameter choice.

3) *Cluster Count (NCLT)*: Table IV-E3 investigates the number of visual prototypes per class:

NCLT	Accuracy	Precision	Recall	F1
8	71.21%	70.52%	69.56%	69.25%
16 (optimal)	72.36%	72.03%	71.04%	70.54%
24	71.59%	71.60%	69.47%	69.60%
32	71.21%	72.12%	70.63%	70.57%

NCLT=16 provides optimal intra-class diversity capture. Both undersampling (8) and oversampling (32) degrade accuracy, indicating a sweet spot exists where clustering balances diversity representation without overfitting.

4) *Reproducibility (Random Seed)*: Table IV-E4 validates training stability:

Seed	Accuracy	Precision	Recall	F1
1	72.4%	72.0%	71.0%	70.5%
2	72.0%	73.2%	70%	69.6%
3	71.4%	71.7%	69.8%	69.8%
	$71.9 \pm 0.4\%$	$72.3 \pm 0.7\%$	$70.3 \pm 0.6\%$	$70 \pm 0.5\%$

Standard deviation of 0.45% across seeds demonstrates reproducibility and suggests minimal variance from stochastic initialization, a desirable property for production deployment.

F. Cross-Dataset Generalization

Table IV-F evaluates transfer across datasets:

Dataset	Classes	Accuracy	Precision	F1
PlantDoc	27	72.36%	72.03%	70.54%
PlantVillage	39	68.9%	68.7%	67.5%
PlantWild	115	61.2%	61.8%	58.9%

Performance decreases with dataset complexity (more classes) and distribution shift (field vs. lab images), as expected. However, absolute accuracy remains competitive without any fine-tuning, demonstrating generalization capacity of frozen CLIP features.

V. DISCUSSION

A. Key Findings

Our results demonstrate that combining frozen vision-language models with lightweight learnable prototypes provides an effective approach for plant disease classification:

- 1) **Strong baseline performance**: 72.36% on PlantDoc represents a 34.1 percentage point improvement over zero-shot CLIP, achieved with only 500K trainable parameters.
- 2) **Data efficiency**: 59.7% accuracy at 20-shot (82.5% relative efficiency) suggests practical deployment in data-limited agricultural settings.
- 3) **Hyperparameter robustness**: Ablation studies show stability across reasonable hyperparameter ranges, facilitating transfer to new datasets.
- 4) **Reproducibility**: Low variance across random seeds ($\pm 0.45\%$) indicates reliable training dynamics.

B. Computational Efficiency

Training on RTX 3060: Baseline 45 minutes, Few-shot (1-shot) 8 minutes, Full ablation suite 6 hours. Prototype caching reduces redundant computation by $4\times$, enabling efficient experimental iteration.

C. Limitations and Honest Reporting

We acknowledge several limitations:

- 1) **Open-set recognition**: AUROC ≈ 0.5 for out-of-distribution detection indicates the frozen adapter architecture was not designed for uncertainty quantification. Specialized OOD detection methods would be necessary for deployment scenarios requiring confidence rejection.
- 2) **Cross-dataset generalization**: 10-15 pp accuracy drops when testing on different datasets without fine-tuning highlight domain shift challenges. Domain adaptation techniques could be explored in future work.
- 3) **Per-class imbalance**: Accuracy ranges 65-85% across disease classes. Rare diseases with fewer training samples perform worse, consistent with standard supervised learning behavior.

- 4) **Scalability to 100+ classes:** Performance on PlantWild (115 classes) drops to 61.2%. Scaling to large taxonomies may require architectural modifications.

D. Ethical Considerations

Accessibility: Open-source implementation and CPU-compatible inference prioritize accessibility for smallholder farmers and extension services in low-resource regions.

Intended Use: MVPDR should be deployed as a *triage and decision-support tool*, not autonomous diagnosis. Human expert validation remains essential, particularly for high-stakes agricultural decisions.

Dataset Bias: Evaluated on three independent benchmarks. Honest reporting of per-class performance variations highlights where model confidence should be lower.

VI. CONCLUSION

MVPDR demonstrates that frozen vision-language models can be effectively adapted for fine-grained agricultural classification through lightweight prototype learning. Our approach achieves strong accuracy (72.36% on PlantDoc) with minimal trainable parameters and demonstrates promising data efficiency for few-shot scenarios. Comprehensive ablation studies and reproducibility analysis provide confidence in the approach’s robustness.

Future directions include: (1) exploring learnable clustering strategies beyond K-means; (2) investigating domain adaptation for field-to-lab transfer; (3) developing uncertainty quantification for deployment; (4) extending to multi-modal fusion incorporating environmental sensor data; and (5) scaling to larger disease taxonomies with hierarchical classification strategies.

The released code and documentation aim to support agricultural AI adoption in resource-limited settings where accessibility and efficiency are paramount.

REFERENCES

- [1] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, et al., “Learning transferable visual models from natural language supervision,” in *International conference on machine learning*, PMLR, 2021, pp. 8748–8763.
- [2] D. Singh, N. Jain, P. Jain, P. Kayal, S. Kumawat, and N. Batra, “Plantdoc: A dataset for visual plant disease detection,” in *Proceedings of the conference on computer vision and pattern recognition workshops*, 2020, pp. 249–253.
- [3] D. P. Hughes and M. Salathé, “An open access repository of images on plant health to enable the development of mobile disease diagnostics,” arXiv preprint arXiv:1511.08060, 2015.
- [4] C. Finn, P. Abbeel, and S. Levine, “Model-agnostic meta-learning for fast adaptation of deep networks,” in *International conference on machine learning*, PMLR, 2017, pp. 1126–1135.
- [5] S. P. Lloyd, “Least squares quantization in pcm,” *IEEE transactions on information theory*, vol. 28, no. 2, pp. 129–137, 1982.